

Fine-tuning LLMs

Fine-tuning is the process of adapting a pre-trained language model to a specific task or domain. Instead of training from scratch, fine-tuning starts with a model trained on large corpora and continues training on task-specific data. This approach requires far less data and compute than training from scratch. Common fine-tuning methods include full parameter fine-tuning and parameter-efficient techniques like LoRA and QLoRA. Fine-tuned models can specialize in particular domains like medical, legal, or technical content. However, RAG often provides a more flexible alternative to fine-tuning, as it can incorporate new information without model retraining.